

Video Forgery Detection

Deep Learning DATS 6303 - Fall 2024

Group 6

Individual report

Author: Khush Shah



Table of Contents:

Table of Contents:	2
Introduction	3
Repository Structure	4
Dataset	5
Proof of Concept with Image Dataset.....	5
Models Tried and Architecture	7
1. VGG Model.....	8
2. Basic CNN Model.....	8
Results	10
Learning Curve	11
Conclusion	13
Future Work	14

Introduction

The exponential growth of multimedia content has revolutionized how we communicate and share information. Approximately 3.2 billion images and 720,000 hours of videos are shared daily ([source](#)). However, this surge in content has also given rise to challenges such as misinformation and deepfake manipulation.

Deepfakes—synthetically generated or altered videos and images—pose a significant threat to digital trust. These manipulations have been used for:

- Political Propaganda: Fake videos of politicians making inflammatory statements.
- Financial Fraud: Identity theft through morphed images.
- Misinformation Campaigns: Spreading false narratives with convincing fake visuals.

The societal impact of deepfakes is profound, ranging from eroding public trust to causing financial losses and reputational damage. Detecting these manipulations is crucial to mitigating their harmful effects.

Our project aims to address this issue by developing robust deep-learning models for forgery detection. By leveraging state-of-the-art datasets like CelebDF and implementing advanced architectures such as GRU, VGG, and 3dCNN, we strive to create a reliable system for identifying manipulated multimedia content.

Repository Structure

The repository is structured into different folders based on the experiments conducted:

1. `cnn_basic/`:
 - Contains code for implementing a basic CNN model.
 - Focused on lightweight architectures for efficient image-based forgery detection.
2. `streamlit/`:
 - Includes the Streamlit app for user interaction.
 - Features dropdown menus for model selection, file upload functionality, and preprocessing steps.
3. `vgg/`:
 - Contains code for adapting the VGG model for forgery detection.
 - Includes auxiliary feature integration (e.g., mean RGB values).
4. `video_mask/`:
 - Focuses on preprocessing video data by extracting frames and applying masks.
 - Implements different approaches for data loading and auxiliary feature extraction.
5. `eda.py`:
 - Performs exploratory data analysis (EDA) on the dataset.
 - Calculates mean pixel values for real and fake images.

Each folder represents a milestone in our journey toward building an effective forgery detection system.

Dataset

We used the CelebDF dataset, a large-scale dataset specifically designed for deepfake detection ([source](#)). The dataset is divided into three categories:

1. Real Videos:
 - Authentic videos of celebrities.
2. Fake Videos:
 - Deepfake-generated videos using advanced synthesis techniques.
3. YouTube Videos:
 - Real-world examples of multimedia content.


The CelebDF dataset contains approximately 6,000 videos, with around 1,100 real videos and the remaining being forged videos. The forged videos are generated using advanced synthesis techniques, making them highly realistic and difficult to detect. This diversity and quality make CelebDF a challenging benchmark for forgery detection models.

Proof of Concept with Image Dataset

Before transitioning to the video dataset, we used an image dataset to conduct a proof of concept (POC). This dataset contained approximately 70,000 images, with around 10,000 real images and the remaining being fake images generated through various methods such as:

- GANs (Generative Adversarial Networks): Used to create highly realistic fake images.
- FaceLabs: A tool for generating manipulated facial images.
- FaceApp: A popular app for facial transformations.

The image dataset allowed us to experiment with preprocessing techniques, auxiliary feature extraction (e.g., mean RGB values), and initial model training. These experiments provided valuable insights into handling real vs. fake data and informed our approach to working with the more complex video datasets.



(a)



(b)

Fig 1a,b : Real and fake image.

Figures 1a and 1b depict the real and distorted images through the faceapp method.

Models Tried and Architecture




Figure 2: Flow Architecture

Figure 2 represents the comprehensive workflow of the forgery detection system, highlighting the sequential steps involved in detecting manipulated multimedia content. The Input Layer consists of a video dataset categorized into real, fake, and YouTube real videos. These datasets serve as the foundation for training and testing the models. The Preprocessing Layer involves converting videos into frames at specific rates (frames per second), followed by augmentation, resizing, and splitting the data into training and testing sets. This ensures that the data is well-prepared for effective model training.

The Model Layer is where the core of the system resides. It includes training deep learning models on preprocessed data and evaluating their performance using unseen test data to ensure generalization. The trained model learns to distinguish between real and fake content based on spatial and temporal features extracted from the frames. Finally, the Output Layer is responsible for detecting forgery in multimedia content based on the predictions made by the model. This layer outputs whether a given video is real or fake, providing a robust solution to combat misinformation caused by deepfakes.

This structured pipeline ensures that each layer contributes to building an efficient and accurate forgery detection system, capable of handling diverse datasets with high-quality synthesis techniques.

1. VGG Model

The VGG (Visual Geometry Group) model is a convolutional neural network (CNN) architecture widely used in image classification tasks. Its simplicity and effectiveness make it a popular choice for deep learning projects.

Key Features:

- Input images were resized to (224x224) dimensions.
- The architecture consisted of multiple convolutional layers with small (3x3) filters followed by max-pooling layers.
- Fully connected layers at the end were modified to include auxiliary data (e.g., mean RGB values of real and fake images) as additional inputs.
- The model was trained using focal loss to handle class imbalance effectively.

Challenges:

- While the VGG model performed well on image datasets, it struggled with temporal features in video datasets due to its lack of sequential processing capabilities.
- High computational requirements made training time-consuming.

2. Basic CNN Model

To reduce resource usage and simplify training, we implemented a basic CNN model with auxiliary features. This model was designed to process images efficiently while incorporating auxiliary data.

Key Features:

- A lightweight architecture with two convolutional layers followed by max-pooling layers.
- Auxiliary data (mean RGB values) was concatenated with flattened image features before passing through fully connected layers.
- The model was trained using focal loss, which helped improve performance on imbalanced datasets.

Results:

- The basic CNN model demonstrated promising results on image datasets but required further optimization for video datasets due to their temporal complexity.

Learning Curve

Our journey involved several milestones that shaped our understanding of forgery detection challenges. We began by exploring repositories like Awesome Deepfakes Detection to identify suitable datasets for our project. After evaluating multiple options, we selected the CelebDF dataset due to its diversity and relevance for deepfake detection.

As a proof of concept (POC), we started with image datasets to understand how to load, preprocess, and segment data into real and fake classes. Preprocessing steps included resizing images, normalizing pixel values, and extracting auxiliary features such as mean RGB values. When these auxiliary features were concatenated with image features at the fully connected layer stage in both VGG and CNN models, the model's accuracy doubled, demonstrating the importance of incorporating additional contextual information.

Exploratory Data Analysis (EDA) revealed key differences between real and fake images. For example, real images had lower red channel values compared to fake images. These insights guided our preprocessing strategies and helped refine our models.

After achieving satisfactory results on image datasets, we transitioned to video datasets. This phase involved extracting frames from videos at specific intervals (e.g., 100 FPS), calculating auxiliary features for each frame, and training models on these frames. The inclusion of auxiliary data significantly improved the model's ability to detect forged content.

We also explored Error Level Analysis (ELA) as a forensic technique for detecting image manipulation (source). ELA analyzes compression artifacts in JPEG images by re-saving them at a lower quality and comparing error levels between the original and recompressed versions. Regions with higher error levels often indicate tampering or editing. Although primarily used for static images, ELA provided valuable insights into preprocessing techniques for identifying manipulated content.

Additionally, we reviewed several research papers to refine our approach:

- "FaceForensics++: Learning to Detect Manipulated Facial Images": This paper introduced a large-scale dataset for detecting facial manipulations and explored deep learning architectures for forgery detection.
- "DeepFake Detection Using Recurrent Neural Networks": The authors proposed using RNNs to analyze temporal inconsistencies in video frames for detecting deepfakes.

- "Exposing DeepFake Videos by Detecting Face Warping Artifacts"): This study focused on identifying warping artifacts introduced during face-swapping processes in deepfake videos.

These papers provided critical insights into state-of-the-art methods for forgery detection and inspired us to integrate auxiliary features into our models. Through these milestones, we gained valuable insights into forgery detection challenges and developed robust preprocessing pipelines tailored to our dataset.

I utilized GenAI, and the code documentation of the used libraries to write down the code. It was a balance of 50-50% for both of them. Reading documentation has helped in the nuances of the improvement of the codes.

Conclusion

This project represents a significant step toward combating misinformation caused by deepfakes and manipulated media. By leveraging state-of-the-art datasets like CelebDF and implementing advanced deep learning architectures (e.g., VGG, CNN), we demonstrated promising results in detecting forgery in multimedia content.

Key Takeaways:

1. Understanding the dataset is critical before implementation—EDA revealed valuable insights that guided our preprocessing strategies.
2. Auxiliary features such as mean RGB values can enhance model performance by providing additional context about real vs. fake data.
3. Lightweight architectures like basic CNNs offer a good balance between performance and resource efficiency for image-based forgery detection tasks.

While our models showed promising results on image datasets, further optimization is required for video datasets due to their temporal complexity.

Future Work

To enhance our forgery detection system further, we propose the following directions:

1. Face Mesh Analysis:
 - Generate 3D face meshes for detected faces in videos.
 - Analyze changes in node values on the mesh to identify subtle manipulations.
2. Image Masking Techniques:
 - Implement masking algorithms to highlight differences between real and altered images.
 - Use these masks as auxiliary features during training for improved accuracy.
3. Multi-modal Approaches:
 - Combine video and audio streams for forgery detection.
 - Analyze inconsistencies between visual (video) and auditory (audio) modalities to detect manipulation.
4. Temporal Feature Extraction:
 - Explore architectures like LSTMs or Transformers for capturing temporal dependencies in video datasets.
 - Train models on sequences of frames instead of individual frames for better context understanding.
5. Real-world Testing:
 - Test models on real-world examples from platforms like YouTube or social media.
 - Evaluate performance on diverse scenarios beyond curated datasets like CelebDF.

By addressing these areas, we aim to build a comprehensive forgery detection system capable of tackling real-world challenges effectively. This extended documentation provides an in-depth overview of your project's progress, achievements, challenges faced, key insights gained through research integration, and future directions! Let me know if you need further refinements!